

HOWL — in brief

Design summary · the full specification runs to 62 pages

What HOWL is

HOWL is a sound-and-complete, consequence-based description-logic reasoner, sibling to GROWL. GROWL **materializes** an OWL 2 RL ontology; HOWL **classifies** a description logic — it computes the complete subsumption hierarchy and coherence of a TBox, including entailments that forward-chaining materialization structurally cannot see.

The two engines answer different questions over the same ontologies:

	GROWL	HOWL
Service	Materialization — forward-chain sound consequences	Classification — complete hierarchy + coherence
Method	Semi-naïve fixpoint over RDF triples	Consequence-based saturation over derived axioms
Logic	OWL 2 RL (a rule fragment)	A named description logic ($EL \rightarrow OWL\ 2\ EL \rightarrow SRIQ$)
Completeness	Sound; deliberately incomplete for OWL-DL	Sound and complete for its fragment

RL cannot introduce existential successors, so it silently misses subsumptions that run *through* an existential — exactly the shape a genus-differentia definition takes. That gap is HOWL's reason to exist.

Why now

The execution trigger has fired, and not from where this document originally expected. MOOSE's Living Ontology validates arbitrary user ontologies through a frozen `TBoxReasoner` port, and ships an honest bounded **stub** in place of a real engine. Two ceilings follow, removable by no other means:

- **Instance-bearing changes cannot pass.** The stub reports every assertion-bearing delta inconclusive, restricting the promotion path to annotation-only edits.
- **Dual-engine attestation cannot exist.** The agreement stage needs a second real engine to cross-check against.

This is a stronger trigger than any originally listed, because the cost of not building is being paid continuously by a shipped system rather than deferred until a pipeline stage matures.

The v0 language

v0 is complete for exactly this, and no more:

$C, D ::= \top \mid \perp \mid A \mid C \sqcap D \mid \exists r.C$
TBox $C \sqsubseteq D \mid C \equiv D \mid \text{DisjointClasses}(C_1 \dots C_n)$
RBox $r \sqsubseteq s \mid r_1 \circ \dots \circ r_n \sqsubseteq t \ (n \geq 2) \mid \text{domain}(r) \sqsubseteq C \mid \text{range}(r) \sqsubseteq C$
ABox $C(a) \mid r(a,b)$

Named: $ELH_{\perp}^{R+}$ **with domain and range**. It is a strict subset of $EL++$ and — subject to two RBox conditions — of the OWL 2 EL profile. The gap to $EL++$ is *not* just nominals: concrete domains, `ObjectHasSelf`, reflexive role inclusions, and the two built-in object properties are all out too.

Everything else is enumerated, never dropped. A construct outside the language goes into `coverage.omitted` with a reference to the offending axiom. A dropped axiom is an incoherence the reasoner subsequently cannot see — the fail-silent hazard this whole design is arranged against.


Three dispositions inside the gate, plus one settled before it:

Disposition	Examples	Behaviour
In-profile	the grammar above	classified
Inert	annotation assertions with a declared annotation property	ignored
Out-of-profile	unions, cardinalities, nominals, datatypes, inverses	enumerated
Consumed (<i>pre-gate</i>)	declarations	no axiom, but read — supplies entity typing and the input signature

Getting the inert/out-of-profile boundary backwards is fatal in either direction: permissive drops logical axioms silently, strict returns *inconclusive* on every documented ontology.

The verdict rule

This is the single most load-bearing invariant in the design. Everything else — the port's fault precedence, the CLI exit codes — is this rule in another dialect.



A run is **complete** exactly when `coverage.omitted` is empty *and* the saturation reached a fixpoint. Two consequences, pulling in opposite directions, and both are the point:

- **A finding is definitive even under partial coverage.** Inference is monotone, so restoring an omitted axiom can only *add* conclusions. A $\perp$ derived from part of the ontology is a $\perp$ in all of it. Reporting that as inconclusive would discard a sound result.
- **The absence of a finding means nothing unless the run was complete.** This is why exit 2 exists, and why a CI configuration treating 2 as success reintroduces exactly the hazard the status discipline was built to close.

Status is therefore **three independent axes**, not one enum — a run can be both out-of-profile *and* budget-capped, and collapsing them forces the report to lie about one:

- `coverage` — which axioms were omitted, and why
- `termination` — fixpoint, or resource-limit


- `findings` — what was actually derived

Cancellation is not on this scale at all: it yields a fault with **no result**, because a cancelled partial state is not reproducible and must not be serialized or hashed.

How it works

The front end

Stage order is load-bearing, and two of its constraints pull against each other: the RBox gates must see property chains *before* decomposition, but must also compare *class expressions*, which only exist after RDF is decoded. Decoding first resolves both — on the observation that **decoding an axiom is not accepting it**.




Three details that are easy to get wrong and silent when you do:

- **Identity is tagged, not bare.** OWL permits punning, so `class(:x)`, `individual(:x)` and `fresh(n)` are disjoint by construction. Without that, a class $:x \sqsubseteq \perp$ and an unrelated individual $:x$ merge and the encoding manufactures a **false inconsistency**.
- **Contexts come from the input signature, not from the axioms.** A class that is declared and otherwise unmentioned still gets a context — otherwise the reasoner silently does not know it exists.
- **Normalization is polarity-sensitive.** A subexpression replaced by fresh X needs $E \sqsubseteq X$ in negative position and $X \sqsubseteq E$ in positive. Get it wrong and `Mother  $\sqsubseteq$  Parent` — the single test that defines the project — is never derived.

Saturation

The calculus is CR1–CR7 in the CEL formulation (HOWL's local numbering; the published paper numbers role hierarchy and composition CR10/CR11). Ranges add no rule — they are removed by the published elimination construction during normalization.

State is **context-owned**: one context per concept, holding its own subsumers, successors, predecessors and inbox. A worker mutates only the context it holds; everything else leaves as a message. Global $S(\cdot) / R(\cdot)$ maps would make every rule firing a write to shared state, which is what the decomposition exists to avoid.



Every clause of that loop is load-bearing:

- **The snapshot includes Δ_n .** If both premises of a two-premise rule arrive in the same round, scanning S_{n-1} means neither sees the other and neither is in the next delta — the conclusion is never derived, by any handler, ever.
- **Novel facts only.** Without subtracting the accumulated store, $A \sqsubseteq B$, $B \sqsubseteq A$ re-derives two stored facts every round forever, and an ordinary cyclic hierarchy burns its budget and returns inconclusive.
- **Every rule fires on every premise.** A multi-premise rule needs one handler per premise; a single handler loses every conclusion whose second premise arrives later, and *which* order occurs is a scheduling artifact.
- **Rounds are deterministic.** `iteration` counts completed global rounds with a barrier, so `max-iterations` cuts at the same point at any worker count. Confluence already covers the complete case; the exposure is entirely in the non-complete outcomes, whose partial state would otherwise be a function of thread timing.

Verification, honestly

The dividing line is **structural vs. model-theoretic**, not soundness vs. completeness. A weakest-precondition checker reasons about program states: it can prove the code implements the calculus it claims to, and cannot prove the calculus's conclusions are entailed.

Property	Z3-reachable?
Per-rule faithfulness — only licensed conclusions are emitted, and every licensed one is	Yes
Dispatch coverage — every rule invoked against every axiom, in every trigger direction	Yes
Initialization — empty store, seeds in the queue , every context active	Yes
Termination (v0)	Yes
Per-rule <i>soundness</i> — each conclusion is entailed	No — proof-inherited
Range-elimination preservation; ABox reduction correctness	No
Global completeness	No — meta-theoretic


All three structural obligations are needed together. Drop any one and a **disabled engine passes verification**: with only the soundness direction, rules returning the empty list satisfy everything; without the initialization obligation, seeds written to the store leave the frontier empty and the driver reports `fixpoint` before firing a rule.

So the honest claim is *verified faithful to a proven calculus* — not "verified sound", not "verified complete".

The ABox obligation is discharged. The direct-edge ABox encoding is HOWL's own construction, not lifted from a published proof. SPEC §5.3 proves it sound and complete by reduction to the published EL++ calculus on a nominal-free CBox and its range-restriction extension, with the ABox's nominals discharged by a canonical-model argument (the published completeness for nominals fails, KKS12). It was reviewed 2026-09-30, so ABox support is inside the sound-and-complete v0 claim.

Integration

HOWL is the unbuilt engine behind MOOSE's `TBoxReasoner`. **The port, not the CLI, is the deliverable.**



Note this is stronger than one-directional composition: the two engines are also **independent oracles checked against each other in production**, where disagreement is an engine fault, never resolved by preferring one.

Four port amendments are prerequisites


`TBoxInput` is `{ triples: Vec<Triple> }` — that is the whole of it. Four capabilities this design relies on cannot cross it as it stands:

	Amendment	Without it
A1	per-document provenance	axiom identity and blank-node separation are lost when documents flatten
A2	import closure attestation	a dangling import yields a false coherent
A3	baseline channel + <code>Option</code> delta	no baseline can be identified; <code>classify</code> is pure
A4	probe-set entailment readback	the agreement stage fails closed — landing HOWL makes the pipeline <i>less</i> usable

An unlanded amendment means the capability is not delivered — never approximated. The adapter never fabricates a value the interface cannot carry.

Two behavioural notes for the swap: HOWL registers as `BoundedDL`, not `ElPlusPlus`, because claiming the latter while implementing a subset is the overclaim this project keeps guarding against. And two of the stub's checks are **deliberately not preserved** — it reports `consistent = false` for subclass cycles and for properties with disjoint domains, neither of which is a DL contradiction. They become modelling lints; a verdict that used to Fail will now Pass with a lint.

Milestones



M1 is where HOWL becomes a DL reasoner, and its acceptance is specific: derives `Mother ⊆ Parent` from the litmus; matches a capability-probed complete oracle as *entailments* rather than serialized taxonomies; meets a pinned benchmark protocol; **never reports coherent from an incomplete run** while still reporting incoherent whenever a contradiction was proven; produces byte-identical reports across worker counts; and discharges the ABox proof obligation.

M5 is where the first consumer's coverage gap actually closes — 33 `owl:inverseOf` uses in its corpus stay out-of-profile until `sriq` lands. `v0` buys instance-bearing verdicts and range coherence, not a clean pass on that consumer's own theory.

What this design refuses to do

- **Guess.** Every unsupported construct is enumerated. Coverage travels on a field the consumer already has to read, never in prose telling someone to check something.
- **Approximate a missing capability.** No baseline channel means no delta — not a delta computed against nothing.
- **Claim more than the proof covers.** Gates match the cited condition exactly; a relaxation that looks semantically fine still needs its own proof, because completeness rests on implementing a calculus someone has proved complete.
- **Report success it did not establish.** The three ways to reach *inconclusive* — out-of-profile, budget, cancellation — are structurally unable to become *coherent*.

The recurring failure mode this document is arranged against is a mechanism that fails silently in the direction that looks like success. Most of the design's apparent over-caution is one instance or another of closing that door.